

HARISH KRISHNAN

Data Engineer

harishk39024@gmail.com | +91 8296847111 | Bengaluru, India | linkedin.com/in/harish-k-3b98b9317 | github.com/Hurry-sh

OBJECTIVE

Computer Science Engineering student and Data Engineering Intern at BigBasket with hands-on experience in AWS data lake pipelines, CDC ingestion, dbt-based data modeling, Airflow-scheduled workflows, machine learning, and data analytics. Seeking an entry-level Data Engineer role focused on building reliable batch and CDC pipelines, SQL-driven data models, and analytics-ready data marts.

SKILLS

Programming Languages: Python, SQL

Data Engineering: PySpark, Apache Airflow, DBT, Trino, Apache Iceberg, Parquet, ETL/ELT

Cloud & Databases: AWS S3, AWS Athena, Amazon Redshift, MySQL

Machine Learning: TensorFlow

Tools / Platforms: Git, VS Code, OpenMetadata, OpenAI Codex, Claude, MCP Servers

WORK EXPERIENCE

Data Engineering Intern

BigBasket

Jan 2026 – Present

Bengaluru, India

- Reduced monthly platform cost from \$35K to \$8K by contributing to BigBasket's migration from Amazon Redshift to an AWS data lake using S3, Apache Iceberg, DBT, Trino, and Athena, while supporting pipelines with a 99.8% success rate.
- Facilitated the migration of essential Change Data Capture (CDC) pipelines that handles millions of rows daily across diverse datasets, including Planogram, E-Retail, BBData Daily and Last-Mile Delivery Pipeline which handles 110+ columns.
- These tables are responsible to handle warehouse operations, useful for Out-Of-Stock RCA, comprises data of millions of users and is useful for multiple downstream analytics, machine learning tasks and reporting.
- Completed multiple AWS Database Migration Service ingestion tasks and Spark jobs to transfer 3B+ rows from source systems into the Bronze layer of the AWS data lake.
- Developed and maintained 100+ Bronze and Silver data assets, along with a Gold-layer finance data mart for monthly opening and closing balances of roughly 2.5M BBDaily users.
- Used Apache Airflow to schedule, monitor, and rerun production data workflows, handling daily processing, backfills, logs, retries, and validation checks.
- Leveraged OpenMetadata to capture table lineage, validate freshness and completeness of created data assets, and enable easier data discovery during development of the Last-Mile Delivery pipeline.

EDUCATION

B.Tech in Computer Science Engineering | CGPA: 8.68

PES University

Aug 2022 – Jul 2026

Bengaluru, India

12th Grade, PCMC | Percentage: 92.8%

Sri Chaitanya PU College

2020 – 2022

Bengaluru, India

10th Grade, CBSE | Percentage: 90.6%

BGS National Public School

2016 – 2020

Bengaluru, India

ACHIEVEMENTS

- 6x Distinction Scholar at PES University.
- Received Academic Excellence Trophy in 12th grade.

PROJECT EXPERIENCE

Redshift to AWS Data Lake Migration

AWS S3, Apache Iceberg, DBT, Trino, Athena, Airflow

- Contributed to migration design for moving Redshift and RDS-sourced datasets into an S3-backed data lake with Bronze, Silver, and Gold layers.
- Developed and tested 100+ DBT models, including incremental logic, schema mappings, and table-level validation queries before exposing migrated tables to downstream users.

- Created validation checks for row counts, duplicates, null values, data types, and partitions to identify migration issues before production usage.

Android Malware Detection using ML and Hybrid Analysis Python, CNN, Grad-CAM, MobSF, QUARK Engine

- Built an Android malware detection workflow in Python using CNNs, Grad-CAM, API call graphs, MobSF, and QUARK Engine for APK classification and hybrid malware analysis.
- Compared the project approach with MAPAS literature benchmark results, documenting reported accuracy and memory trade-offs against existing Android malware detection techniques.
- Referenced MAPAS benchmark results showing approximately 99% accuracy and 2.2 GB RAM usage, compared with MaMaDroid's reported 34 GB RAM usage.
- Co-authored "Detection of Mobile Malware (ANDROID) using ML and Hybrid Analysis," published in ACM International Conference Proceedings and indexed in Ei Compendex and Scopus.
- Paper link: dl.acm.org/doi/10.1145/3726101.3726103

LINKS

Portfolio: hurry-sh.github.io/Portfolio/

GitHub: github.com/Hurry-sh

LinkedIn: linkedin.com/in/harish-k-3b98b9317